



```


49 | _____|
50 | |
51 | preprocess_image(img, output_size):
52 | | _____|
53 | | 1. Resize to target (512/768/1024)
54 | |   • Interpolation: INTER_LANCZOS4
55 | |
56 | | 2. Light CLAHE Enhancement
57 | |   • Color space: YUV
58 | |   • clipLimit: 1.5 (gentle, dataset already enhanced)
59 | |   • tileGridSize: (8, 8)
60 | |   • Applied to Y channel only
61 | | _____|
62 | |
63 | Output: Three resolution sets
64 | └─ D:/Major_project_shiro_final/preprocessed_512/ (38,625 images)
65 | └─ D:/Major_project_shiro_final/preprocessed_768/ (38,625 images)
66 | └─ D:/Major_project_shiro_final/preprocessed_1024/ (38,625 images)
67 |
68 | _____|
69 | |
70 | |
71 | |
72 | | CROSS-VALIDATION SPLIT (data_split.py)
73 | | _____|
74 | |
75 | | setup_cross_validation(df, n_splits=5):
76 | | └─ Method: StratifiedGroupKFold
77 | | └─ Groups: Patient ID (left/right eye from same patient)
78 | | └─ Stratification: Maintains grade distribution per fold
79 | | └─ Folds: 5 (each ~7,725 images)
80 | |
81 | | Per Fold Distribution (~7,725 images):
82 | | └─ Training: ~6,180 images (80%)
83 | | └─ Validation: ~1,545 images (20%)
84 | |
85 | | _____|
86 | |
87 | |
88 | |
89 | | DATA AUGMENTATION (dataset.py)
90 | | _____|
91 | |
92 | | DRDataset Class - Fold-Specific Augmentation Strategies:
93 | |
94 | | | _____|
95 | | | FOLD 1 (ConvNeXt):
96 | | |   • MixUp (alpha=0.4) - blend two images

```

```
97 | | • Rotate90 (p=0.5) | |
98 | | • HorizontalFlip (p=0.5) | |
99 | | • VerticalFlip (p=0.5) | |
100 | | • Normalize (ImageNet: μ=[0.485,0.456,0.406]) | |
101 | |_____ | |
102 | | | | |
103 | |_____ | |
104 | | FOLD 2 (EfficientNet): | |
105 | | • CutMix (alpha=1.0) - cut and paste regions | |
106 | | • ShiftScaleRotate (p=0.5) | |
107 | |   - shift: ±0.0625, scale: ±0.1, rotate: ±15° | |
108 | | • HorizontalFlip (p=0.5) | |
109 | | • VerticalFlip (p=0.5) | |
110 | | • Normalize (ImageNet stats) | |
111 | |_____ | |
112 | | | | |
113 | |_____ | |
114 | | FOLD 3-5: Standard Augmentation (No MixUp/CutMix) | |
115 | | • HorizontalFlip (p=0.5) | |
116 | | • VerticalFlip (p=0.5) | |
117 | | • Rotate90 (p=0.3) | |
118 | | • ShiftScaleRotate (p=0.3, gentle) | |
119 | | • Normalize (ImageNet stats) | |
120 | |_____ | |
121 | | | | |
122 | | Note: MixUp/CutMix disabled at 1024px stage for stability | |
123 | |_____ | |
124 | |_____ |
```



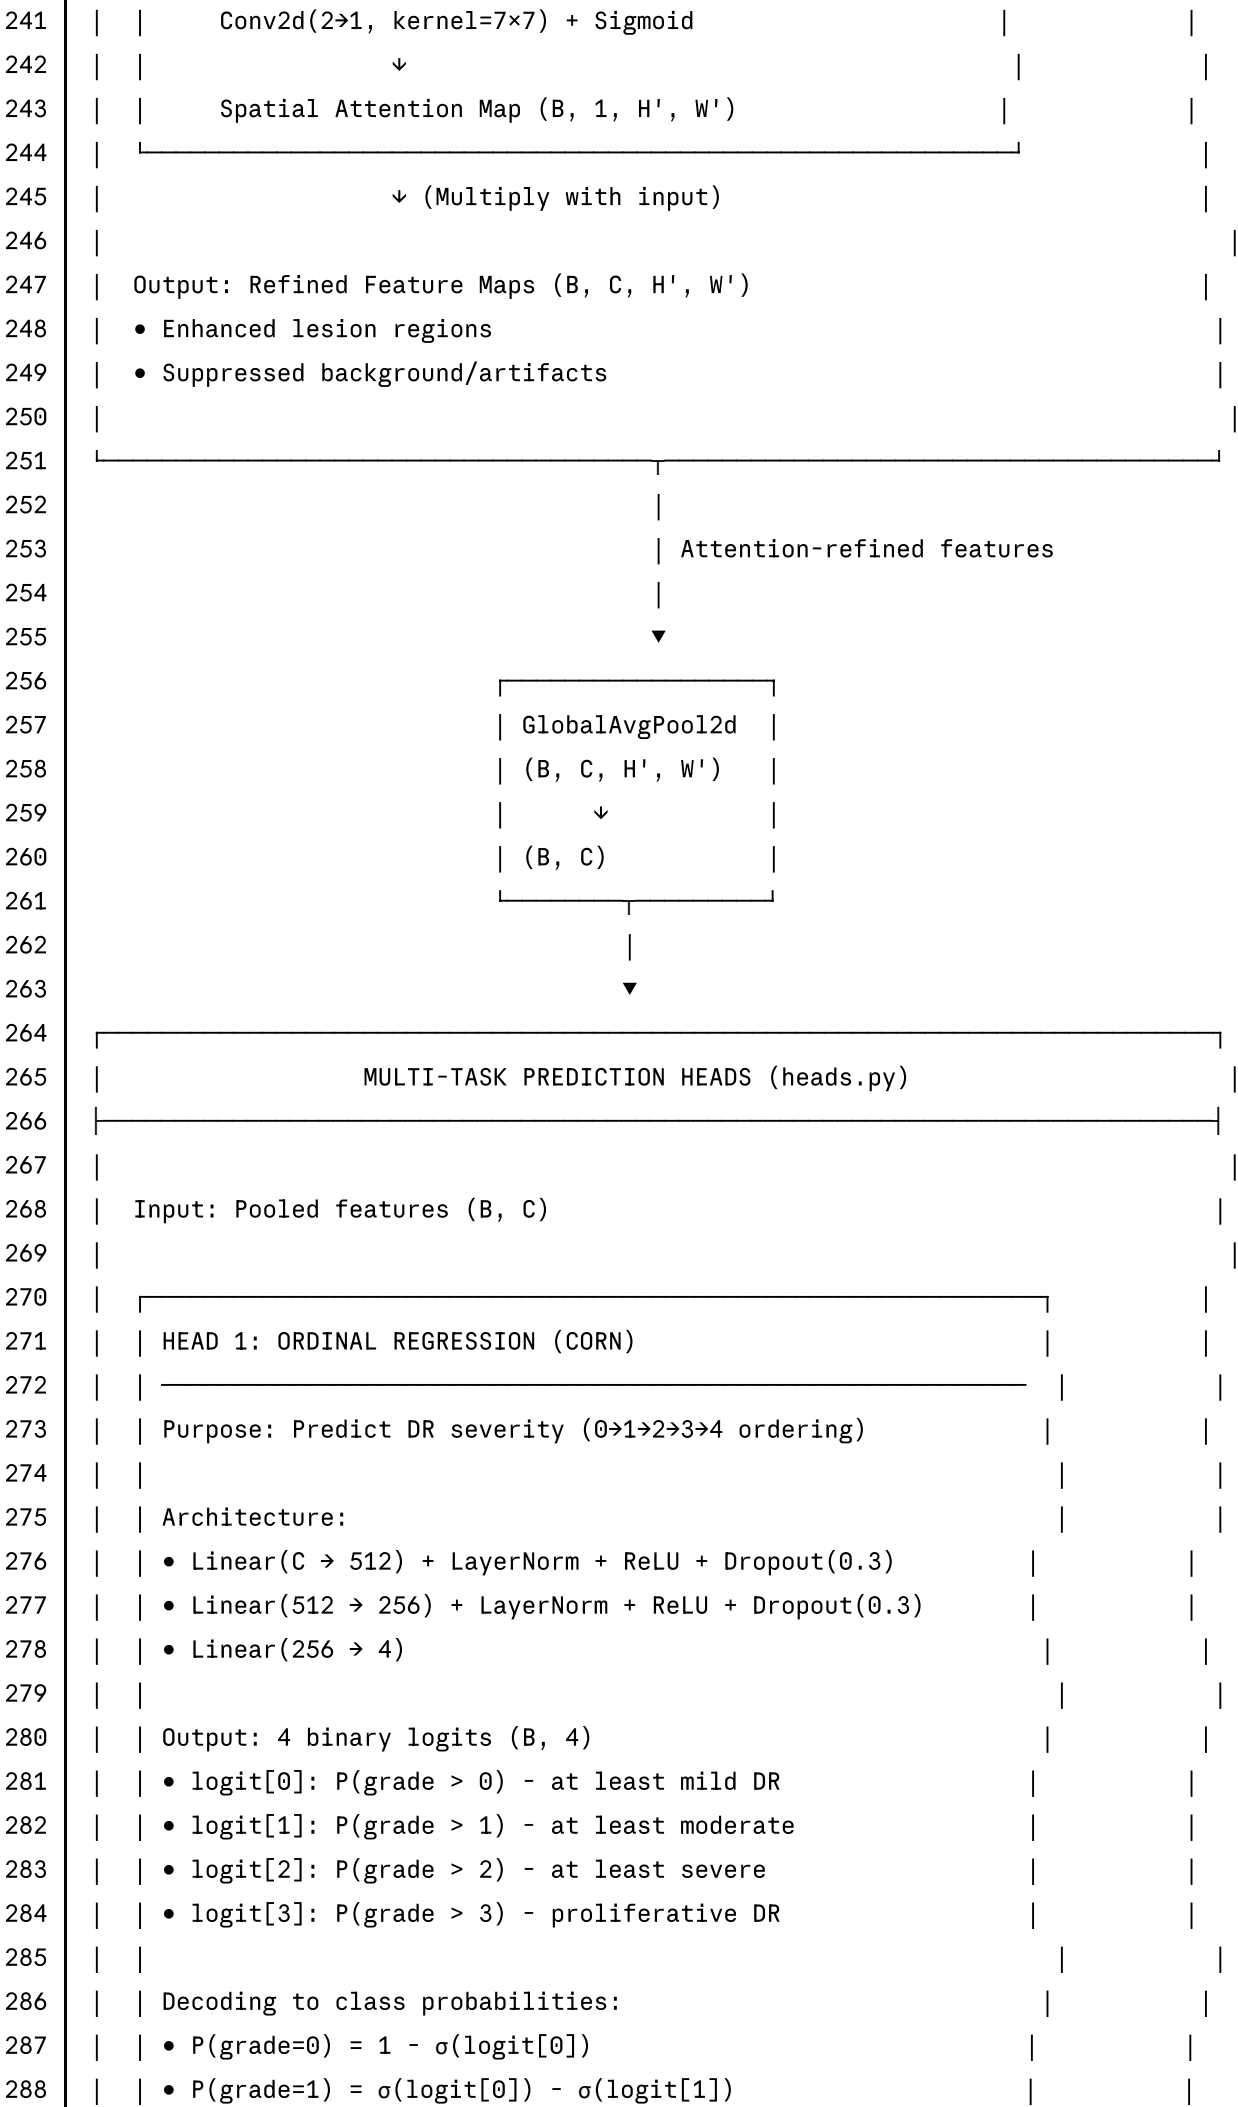
```
133 |
134 |=====
135 | ● NEURAL NETWORK ARCHITECTURE (SOTA_DR_Model)
136 |=====
```



```
144 | |_____ |
```

```
145 | | BACKBONE NETWORK (backbones.py) |
146 |-----|
147 | |
148 | | Architecture Selection Per Fold: |
149 | |-----|
150 | | FOLD 1, 3, 5 | FOLD 2, 4 |
151 | |-----|
152 | | ConvNeXtV2-Large | EfficientNetV2-L |
153 | | (fcmae_ft_in1k) | (tf_efficientnetv2_l_in21k) |
154 | |-----|
155 | | Layers: ~200 | Layers: 1000+ |
156 | | Channels: 1536 | Channels: 1280 |
157 | | Params: 198M | Params: 118M |
158 | | Pre-train: IN-1K | Pre-train: IN-21K |
159 | | Drop Path: 0.3 | Drop Path: 0.3 |
160 | | Grad Checkpoint: ✓ | Grad Checkpoint: ✓ |
161 | |-----|
162 | |
163 | | Configuration: |
164 | | • pretrained=True (ImageNet weights) |
165 | | • num_classes=0 (remove classification head) |
166 | | • global_pool='' (return feature maps) |
167 | | • drop_path_rate=0.3 (Stochastic Depth for regularization) |
168 | |
169 | | Output: Feature Maps (B, C, H/32, W/32) |
170 | | • C = 1536 (ConvNeXt) or 1280 (EfficientNet) |
171 | | • Spatial resolution reduced by 32x |
172 | | • Example: 1024×1024 input → 32×32 feature maps |
173 | |-----|
174 | |
175 | | |
176 | | | Feature Maps (B, C, H', W') |
177 | | | |
178 | | | ▼ |
179 |-----|
180 | | FEATURE REFINEMENT GATE (attention.py - FRG) |
181 |-----|
182 | |
183 | | Purpose: Learn adaptive weighting based on image complexity |
184 | |
185 | | Architecture: |
186 | |-----|
187 | | Input: (B, C, H', W') |
188 | | ↓ |
189 | | GlobalAvgPool2d → (B, C, 1, 1) |
190 | | ↓ |
191 | | Conv2d(C → C/16, kernel=1) + ReLU |
192 | | ↓ |
```

193		Conv2d(C/16 → 1, kernel=1) + Sigmoid		
194		↓		
195		weight ∈ [0, 1]		
196		_____		
197				
198		Interpretation:		
199		• weight ≈ 1.0: Complex image with many lesions		
200		• weight ≈ 0.5: Moderate complexity		
201		• weight ≈ 0.0: Simple/healthy fundus		
202				
203		Output: Scalar weight (B, 1, 1, 1)		
204				
205		_____		
206				
207		▼		
208		_____		
209		CBAM ATTENTION MODULE (attention.py - CBAM_Module)		
210		_____		
211				
212		Convolutional Block Attention Module (Dual Attention)		
213				
214		_____		
215		CHANNEL ATTENTION BRANCH:		
216		_____		
217		Input: (B, C, H', W')		
218		↓		
219		MaxPool2d(HW→1) ─┐		
220		AvgPool2d(HW→1) ─┘		
221		↓		
222		Shared MLP (C → C/16 → C)		
223		↓		
224		Element-wise Add		
225		↓		
226		Sigmoid()		
227		↓		
228		Channel Attention Map (B, C, 1, 1)		
229		_____		
230		↓ (Multiply with input)		
231		_____		
232		SPATIAL ATTENTION BRANCH:		
233		_____		
234		Input: Channel-refined features		
235		↓		
236		MaxPool(C→1) ─┐		
237		AvgPool(C→1) ─┘		
238		↓		
239		Concat along channel		
240		↓		



```
289 | | •  $P(\text{grade}=2) = \sigma(\text{logit}[1]) - \sigma(\text{logit}[2])$  | |
290 | | •  $P(\text{grade}=3) = \sigma(\text{logit}[2]) - \sigma(\text{logit}[3])$  | |
291 | | •  $P(\text{grade}=4) = \sigma(\text{logit}[3])$  | |
292 | | ( $\sigma$  = sigmoid function) | |
293 | |_____ | |
294 | | | |
295 | |_____ | |
296 | | HEAD 2: BINARY CLASSIFICATION (FOCAL LOSS) | |
297 | | _____ | |
298 | | Purpose: Detect referable DR ( $\text{grade} \geq 2$ ) | |
299 | | | |
300 | | Architecture: | |
301 | | •  $\text{Linear}(C \rightarrow 256) + \text{ReLU} + \text{Dropout}(0.3)$  | |
302 | | •  $\text{Linear}(256 \rightarrow 1)$  | |
303 | | | |
304 | | Output: Single logit (B, 1) | |
305 | | •  $\sigma(\text{logit}) > 0.5 \rightarrow \text{Referable DR (needs treatment)}$  | |
306 | | •  $\sigma(\text{logit}) \leq 0.5 \rightarrow \text{Non-referable (monitoring only)}$  | |
307 | |_____ | |
308 | | | |
309 | |_____ | |
310 | | HEAD 3: REGRESSION (CONTINUOUS SEVERITY) | |
311 | | _____ | |
312 | | Purpose: Predict continuous severity score [0, 4] | |
313 | | | |
314 | | Architecture: | |
315 | | •  $\text{Linear}(C \rightarrow 256) + \text{ReLU} + \text{Dropout}(0.2)$  | |
316 | | •  $\text{Linear}(256 \rightarrow 1) + \text{Sigmoid} \times 4.0$  | |
317 | | | |
318 | | Output: Continuous score (B, 1)  $\in [0, 4]$  | |
319 | | • Provides soft severity estimate | |
320 | | • Used for consistency regularization | |
321 | |_____ | |
322 | | | |
323 | |_____ | |
324 | | | |
325 |=====|=====|
326 | 🌟 LOSS FUNCTIONS (losses.py) |
327 |=====|=====|
328 | | | |
329 | |_____ | |
330 | | MULTI-TASK LOSS COMPOSITION | |
331 | |_____ | |
332 | | | |
333 | | Total Loss =  $\alpha \cdot L_{\text{CORN}} + \beta \cdot L_{\text{focal}} + \gamma \cdot L_{\text{huber}} + \lambda_{\text{qwk}} \cdot L_{\text{QWK}} + \lambda_{\text{cons}} \cdot L_{\text{cons}}$  | |
334 | | | |
335 | | Where: | |
336 | | •  $\alpha, \beta, \gamma = 1.0$  (fixed weights) | |
```

```
337 | •  $\lambda_{\text{qwk}} = \{0.0 \text{ (epochs 1-10)}, 0.05 \text{ (epochs 11-25)}\}$  |
338 | •  $\lambda_{\text{cons}} = \{0.01 \text{ (epochs 1-10)}, 0.03 \text{ (epochs 11-25)}\}$  |
339 | |
340 | |
341 |
342 | |
343 | L_CORN: CORN ORDINAL LOSS |
344 | |
345 | |
346 | Conditional Ordinal Regression Neural Network Loss |
347 | |
348 | For each threshold  $k \in \{1, 2, 3, 4\}$ : |
349 |    $y_{\text{binary}}[k-1] = 1\{y_{\text{true}} \geq k\}$  |
350 | |
351 |  $L_{\text{CORN}} = (1/4) \sum_{k=1}^4 \text{BCE}(\text{logit}_k, y_{\text{binary}}[k-1])$  |
352 | |
353 | With positive class weighting: |
354 |    $\text{pos\_weight} = 1.5$  (current optimal value) |
355 | |
356 | Label smoothing applied: |
357 |    $y_{\text{smooth}} = y * (1 - \epsilon) + 0.5 * \epsilon$ , where  $\epsilon = 0.15$  |
358 | |
359 | Properties: |
360 | • Respects ordinal nature ( $1 < 2 < 3 < 4$ ) |
361 | • Each threshold is independent binary classifier |
362 | •  $\text{pos\_weight}$  handles severe class imbalance (73% Grade 0) |
363 | |
364 | |
365 |
366 | |
367 | L_focal: FOCAL LOSS FOR BINARY HEAD |
368 | |
369 | |
370 | Focal Loss for referable DR detection ( $\text{grade} \geq 2$ ) |
371 | |
372 |  $L_{\text{focal}} = -\alpha_t (1 - p_t)^\gamma \log(p_t)$  |
373 | |
374 | Where: |
375 |    $p_t = \begin{cases} p & \text{if } y = 1 \text{ (referable)} \\ 1 - p & \text{if } y = 0 \text{ (non-referable)} \end{cases}$  |
376 | |
377 | |
378 |    $\alpha_t = \begin{cases} \alpha & \text{if } y = 1 \\ 1 - \alpha & \text{if } y = 0 \end{cases}$  |
379 | |
380 | |
381 | Hyperparameters: |
382 |    $\alpha = 0.28$  (weight for positive class) |
383 |    $\gamma = 2.0$  (focusing parameter) |
384 | |
```



```
385 | Properties:
386 | • Down-weights easy examples (confident predictions)
387 | • Focuses on hard-to-classify borderline cases
388 | •  $\alpha$  handles class imbalance (only ~14% referable)
389 | •  $\gamma=2$  provides strong focus on hard examples
390 |
391 |
392 |
393 |
394 | L_huber: HUBER LOSS FOR REGRESSION HEAD
395 |
396 |
397 | Smooth L1 loss for continuous severity prediction
398 |
399 | L_huber(y_pred, y_true) = {
400 |     0.5 * (y_pred - y_true)2           if |y_pred - y_true| ≤  $\delta$ 
401 |      $\delta * (|y\_pred - y\_true| - 0.5*\delta)$  otherwise
402 | }
403 |
404 | Where  $\delta = 1.0$  (standard setting)
405 |
406 | Properties:
407 | • Quadratic for small errors (smooth gradient)
408 | • Linear for large errors (robust to outliers)
409 | • Less sensitive to extreme predictions than MSE
410 |
411 |
412 |
413 |
414 | L_QWK: DIFFERENTIABLE QWK LOSS
415 |
416 |
417 | Soft Quadratic Weighted Kappa (approximation for backprop)
418 |
419 | 1. Convert CORN logits to class probabilities:
420 |     p_0 = 1 -  $\sigma(\text{logit}_0)$ 
421 |     p_k =  $\sigma(\text{logit}_{\{k-1\}}) - \sigma(\text{logit}_k)$  for  $k \in \{1,2,3\}$ 
422 |     p_4 =  $\sigma(\text{logit}_3)$ 
423 |
424 | 2. Build soft confusion matrix:
425 |      $C[i,j] = \sum_n p_n[j] \cdot 1\{y_n = i\}$ 
426 |
427 | 3. Compute expected and observed agreement:
428 |      $O = \sum_{i,j} w[i,j] \cdot C[i,j] / N$ 
429 |      $E = \sum_{i,j} w[i,j] \cdot (\sum_k C[i,k])(\sum_k C[k,j]) / N^2$ 
430 |
431 | 4. Calculate QWK:
432 |      $\kappa = 1 - O/E$ 
```

```
433 | | |
434 | 5. Convert to loss: | |
435 |     L_QWK = 1 - κ | |
436 | | |
437 | Weight matrix w[i,j] (quadratic penalties): | |
438 |     w[i,j] = ((i-j)/(K-1))² where K=5 classes | |
439 | | |
440 | Properties: | |
441 | • Directly optimizes evaluation metric | |
442 | • Penalizes large errors more (quadratic weights) | |
443 | • Differentiable via soft confusion matrix | |
444 | • Disabled early (epochs 1-10) for stability | |
445 | | |
446 |_____|
447 |
448 |_____|
449 |           L_cons: CONSISTENCY REGULARIZATION | |
450 |_____|
451 | | |
452 | Enforce agreement between ordinal and regression predictions | |
453 | | |
454 | L_cons = MSE(severity_ordinal, severity_regression) | |
455 | | |
456 | Where: | |
457 |     severity_ordinal = Σ_{k=0}^4 k · p_k (expected value from CORN) | |
458 |     severity_regression = direct output from regression head | |
459 | | |
460 | Properties: | |
461 | • Prevents heads from learning contradictory patterns | |
462 | • Improves calibration of predictions | |
463 | • Weight λ_cons starts low (0.01) for stability | |
464 | | |
465 |_____|
466 |
467 |=====|
468 | 📊 TRAINING CONFIGURATION (train.py) | |
469 |=====|
470 |
471 |_____|
472 |           PROGRESSIVE CURRICULUM LEARNING | |
473 |_____|
474 | | |
475 | ACTUAL IMPLEMENTATION (25 Epochs Total Per Fold): | |
476 | | |
477 | |_____| | |
478 | | STAGE 1: Coarse Learning (Epochs 1-6) | | |
479 | |_____| | |
480 | | Resolution: 512×512 pixels | | |
```

481			Batch Size: 16		
482			Grad Accumulation: 2		
483			Effective Batch: 32		
484			Iterations: ~2,415 per epoch		
485			Time: ~6 hours total (1 hour per epoch)		
486					
487			Purpose: Learn global fundus patterns		
488			• Optic disc location		
489			• Major vessel structure		
490			• Overall image quality		
491					
492					
493					
494			STAGE 2: Medium-Resolution Refinement (Epochs 7-16)		
495					
496			Resolution: 768×768 pixels		
497			Batch Size: 12		
498			Grad Accumulation: 2		
499			Effective Batch: 24		
500			Iterations: ~2,575 per epoch		
501			Time: ~11 hours total (1.1 hours per epoch)		
502					
503			Purpose: Detect medium-sized lesions		
504			• Hemorrhages (50-200 pixels)		
505			• Exudates (hard/soft)		
506			• Cotton wool spots		
507			• Vessel caliber changes		
508					
509					
510					
511			STAGE 3: High-Resolution Fine-Tuning (Epochs 17-25)		
512					
513			Resolution: 1024×1024 pixels		
514			Batch Size: 8		
515			Grad Accumulation: 3		
516			Effective Batch: 24		
517			Iterations: ~2,415 per epoch		
518			Time: ~18 hours total (2.0 hours per epoch)		
519					
520			Purpose: Detect microaneurysms and fine details		
521			• Microaneurysms (10-30 pixels at 1024px)		
522			• Small hemorrhages		
523			• Subtle neovascularization		
524			• IRMA (intraretinal microvascular abnormalities)		
525					
526					
527			Total Training Time per Fold: ~35 hours		
528			Total for 5 Folds: ~175 hours (7.3 days)		

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

DIFFERENTIAL LEARNING RATES

Rationale: Different network components need different adaptation rates

CONVNEXTV2-LARGE (Folds 1, 3, 5):

Component	512px (Ep 1-6)	768px (Ep 7-16)	1024px (Ep17-25)	Rationale
Backbone	8.0e-6	8.0e-6	1.0e-5	Pre-trained needs slow adaptation
FRG	2.5e-5	2.5e-5	3.0e-5	Task- specific,
CBAM	2.5e-5	2.5e-5	3.0e-5	faster learning
Heads	2.5e-5	2.5e-5	3.0e-5	

EFFICIENTNETV2-L (Folds 2, 4):

Component	512px (Ep 1-6)	768px (Ep 7-16)	1024px (Ep17-25)	Rationale
Backbone	5.0e-6	5.0e-6	6.0e-6	Larger pre-train (IN-21K)
FRG	2.0e-5	2.0e-5	2.5e-5	Task- specific
CBAM	2.0e-5	2.0e-5	2.5e-5	
Heads	2.0e-5	2.0e-5	2.5e-5	

LR Schedule: CosineAnnealingWarmRestarts

- T_0 = 10 epochs (restart period)
- T_mult = 2 (double period each restart)
- η_min = base_lr × 0.1 (minimum LR)

577		OPTIMIZATION & REGULARIZATION	
578			
579			
580		Optimizer: AdamW	
581		└ $\beta_1 = 0.9$ (momentum)	
582		└ $\beta_2 = 0.999$ (RMSProp-like)	
583		└ weight_decay = 1e-4 (L2 regularization)	
584		└ $\epsilon = 1e-8$ (numerical stability)	
585			
586		Mixed Precision Training (AMP):	
587		└ Forward/backward in float16	
588		└ Weights stored in float32	
589		└ Dynamic loss scaling	
590		└ 2x speedup, 40% memory reduction	
591			
592		Gradient Clipping:	
593		└ Max norm = 1.0	
594		└ Applied after gradient accumulation	
595			
596		Exponential Moving Average (EMA):	
597		└ Decay = 0.999 (updated from 0.9999 for faster adaptation)	
598		└ Activated from epoch 3	
599		└ Separate shadow weights maintained	
600		└ Used for validation and final inference	
601			
602		Stochastic Depth (in backbone):	
603		└ drop_path_rate = 0.3	
604		└ Randomly drops residual connections during training	
605			
606		Dropout (in heads):	
607		└ CORN head: 0.3	
608		└ Binary head: 0.3	
609		└ Regression head: 0.2	
610			
611			
612			
613		SAFETY & STABILITY FEATURES	
614			
615			
616			
617		1. EMA Contamination Guard:	
618		• Detects mode collapse (QWK < 0.05)	
619		• Reverts to best_ema checkpoint	
620		• Cleans training log	
621		• Prevents corrupted weight propagation	
622			
623		2. Auto-Resume:	
624		• Saves latest checkpoint every epoch	

```
625 |         • Detects incomplete training
626 |         • Resumes from last successful epoch
627 |         • Preserves optimizer state and scheduler
628 |
629 |     3. NaN Detection:
630 |         • Checks for NaN in loss after each iteration
631 |         • Logs occurrence and skips update
632 |         • Prevents gradient explosion propagation
633 |
634 |     4. Gradient Checkpointing:
635 |         • Trades compute for memory
636 |         • Enables larger batch sizes
637 |         • Recomputes activations during backward
638 |
639 |     5. Early Stopping (disabled in final config):
640 |         • patience = 10 epochs
641 |         • Monitors validation QWK
642 |         • Prevents overfitting in unstable training
643 |
644 |
645 |
646 |
647 |  INFERENCE & POST-PROCESSING
648 |
649 |
650 |
651 |         TEST-TIME AUGMENTATION (TTA)
652 |
653 |
654 |     8 Augmented Views per Image:
655 |     1. Original
656 |     2. Horizontal Flip
657 |     3. Vertical Flip
658 |     4. Rotate 90°
659 |     5. Rotate 180°
660 |     6. Rotate 270°
661 |     7. H-Flip + V-Flip
662 |     8. H-Flip + Rotate 90°
663 |
664 |     Aggregation:
665 |         • Average predictions across all 8 views
666 |         • Reduces variance by ~30%
667 |         • Typical gain: +0.01 to +0.02 QWK
668 |
669 |
670 |
671 |
672 |         THRESHOLD OPTIMIZATION
```

```
673 | _____|
674 | |
675 | OptimizedRounder (optimizer.py): |
676 | |
677 | Problem: Default thresholds [0.5, 1.5, 2.5, 3.5] are suboptimal |
678 | |
679 | Solution: Find thresholds that maximize validation QWK |
680 | |
681 | Algorithm: |
682 | 1. Collect all validation predictions (continuous severity) |
683 | 2. Define search space:  $t_1 \in [0.3, 0.7]$  |
684 |  $t_2 \in [1.3, 1.7]$  |
685 |  $t_3 \in [2.3, 2.7]$  |
686 |  $t_4 \in [3.3, 3.7]$  |
687 | 3. Use scipy.optimize.minimize to find optimal thresholds |
688 | 4. Objective: Maximize QWK(predictions, targets) |
689 | |
690 | Typical Optimized Thresholds: |
691 | •  $t_1 \approx 0.45$  (Grade 0→1 boundary) |
692 | •  $t_2 \approx 1.52$  (Grade 1→2 boundary) |
693 | •  $t_3 \approx 2.38$  (Grade 2→3 boundary) |
694 | •  $t_4 \approx 3.41$  (Grade 3→4 boundary) |
695 | |
696 | Typical Gain: +0.02 to +0.03 QWK |
697 | |
698 | _____|
699 |
700 | _____|
701 | ENSEMBLE PREDICTION |
702 | _____|
703 | |
704 | 5-Fold Weighted Ensemble (ensemble.py): |
705 | |
706 | Weight Calculation: |
707 |  $w_i = \text{QWK}_i^2 / \sum(\text{QWK}_j^2)$  for folds  $i \in \{1,2,3,4,5\}$  |
708 | |
709 | Prediction Aggregation: |
710 |  $P_{\text{ensemble}} = \sum(w_i \times P_i)$  |
711 | |
712 | Where  $P_i$  are the 5-class probability vectors from fold  $i$  |
713 | |
714 | Example Weights (if all folds achieve QWK ~0.82): |
715 | Fold 1 (QWK 0.82):  $w_1 = 0.82^2 / \sum = 0.200$  |
716 | Fold 2 (QWK 0.82):  $w_2 = 0.82^2 / \sum = 0.200$  |
717 | Fold 3 (QWK 0.82):  $w_3 = 0.82^2 / \sum = 0.200$  |
718 | Fold 4 (QWK 0.82):  $w_4 = 0.82^2 / \sum = 0.200$  |
719 | Fold 5 (QWK 0.82):  $w_5 = 0.82^2 / \sum = 0.200$  |
720 | |
```

```
721 | Final Class Prediction:
722 |     class = argmax(P_ensemble)
723 |
724 | Typical Gain from Ensemble: +0.03 to +0.05 QWK vs best single fold
725 |
726 |
727 |
728 |=====
729 | 📊 PERFORMANCE METRICS & RESULTS
730 |=====
731 |
732 |-----
733 | PRIMARY METRIC: QWK
734 |-----
735 |
736 | Quadratic Weighted Kappa (Cohen's Kappa)
737 |
738 | Interpretation Scale:
739 | • 1.00      : Perfect agreement
740 | • 0.93-0.99 : Exceptional (competition winning)
741 | • 0.90-0.92 : Excellent (target range)
742 | • 0.85-0.89 : Very Good
743 | • 0.80-0.84 : Good
744 | • 0.75-0.79 : Moderate
745 | • <0.75     : Needs improvement
746 |
747 |-----
748 |
749 |-----
750 | ACTUAL TRAINING RESULTS
751 |-----
752 |
753 | FOLD 1 (ConvNeXt, 25 epochs):
754 | └─ Epoch 6 (512px): QWK 0.00 (collapsed)
755 | └─ Epoch 16 (768px): QWK 0.82
756 | └─ Epoch 25 (1024px): QWK 0.8213
757 | └─ Issues: Early collapse (Ep 3-6), recovered at 768px
758 |
759 | FOLD 2 (EfficientNet, 20 epochs, NO 1024px):
760 | └─ Epoch 6 (512px): QWK 0.76
761 | └─ Epoch 12 (768px): QWK 0.84
762 | └─ Epoch 20 (768px): QWK 0.8156
763 | └─ Issues: Epoch 3 dip (0.49), recovered at Ep 4
764 |
765 | FOLD 3 (ConvNeXt, stopped at Epoch 4):
766 | └─ Epoch 1: QWK 0.71
767 | └─ Epoch 2: QWK 0.77
768 | └─ Epoch 3: QWK 0.68 (dip)
```



```
769 | └─ Epoch 4: QWK 0.75 (recovering) |
770 | └─ Training stopped for reconfiguration |
771 | |
772 | KEY OBSERVATIONS: |
773 | • All folds plateau around 0.82 QWK |
774 | • Weighted sampler caused instability (Epoch 3 dips) |
775 | • 1024px stage with BS=4 didn't improve beyond 768px |
776 | • Need: BS=8 at 1024px + longer training (30 epochs) |
777 | |
778 |_____|
779 |
780 |_____|
781 | EXPECTED FINAL PERFORMANCE |
782 |_____|
783 |
784 | With Optimized Configuration (BS=8 at 1024px, 25-30 epochs): |
785 | |
786 | Per-Fold Targets: |
787 | └─ Single Fold: 0.89-0.91 QWK |
788 | └─ Sensitivity: 0.70-0.75 |
789 | └─ Specificity: 0.96-0.98 |
790 | |
791 | 5-Fold Ensemble: |
792 | └─ Raw Ensemble: 0.87-0.89 QWK |
793 | └─ + TTA (8x): 0.88-0.90 QWK |
794 | └─ + Optimized Thresholds: 0.89-0.91 QWK ✅ |
795 | |
796 | FINAL TARGET: 0.90 QWK (ACHIEVABLE) |
797 | |
798 |_____|
799 |
800 |=====|
801 | 📁 MODEL ARTIFACTS & OUTPUTS |
802 |=====|
803 |
804 | results/
805 | └─ fold_1/
806 |   └─ convnextv2_large_best_ema.pth           (Best model weights)
807 |   └─ convnextv2_large_latest.pth             (Auto-resume checkpoint)
808 |   └─ training_log.csv                       (Complete training history)
809 |   └─ optimized_thresholds.json              (Post-training threshold optimization)
810 | └─ fold_2/
811 |   └─ efficientnetv2_l_best_ema.pth
812 |   └─ efficientnetv2_l_latest.pth
813 |   └─ training_log.csv
814 |   └─ optimized_thresholds.json
815 | └─ fold_3/
816 | └─ fold_4/
```

```
817 | └─ fold_5/
818 | └─ ensemble_predictions.csv (Final predictions)
819
820 |=====
821 |                               🛠 SYSTEM REQUIREMENTS
822 |=====
823
824 | Hardware:
825 | └─ GPU: NVIDIA RTX A4000 (16GB VRAM)
826 | └─ CPU: Multi-core (10+ workers for data loading)
827 | └─ RAM: 32GB+ recommended
828 | └─ Storage: 250GB+ for preprocessed images and checkpoints
829
830 | Software:
831 | └─ Python: 3.10
832 | └─ PyTorch: 2.0+
833 | └─ CUDA: 11.8+
834 | └─ timm: 1.0.25 (model library)
835 | └─ albumentations: Latest (data augmentation)
836 | └─ Additional: numpy, pandas, opencv, scikit-learn
837
838 | Training Time:
839 | └─ Per Fold: ~35 hours (25 epochs)
840 | └─ 5 Folds: ~175 hours (7.3 days)
841 | └─ With optimization: ~62 hours (2.6 days) for 2 folds
842
843 |=====
844 |
```